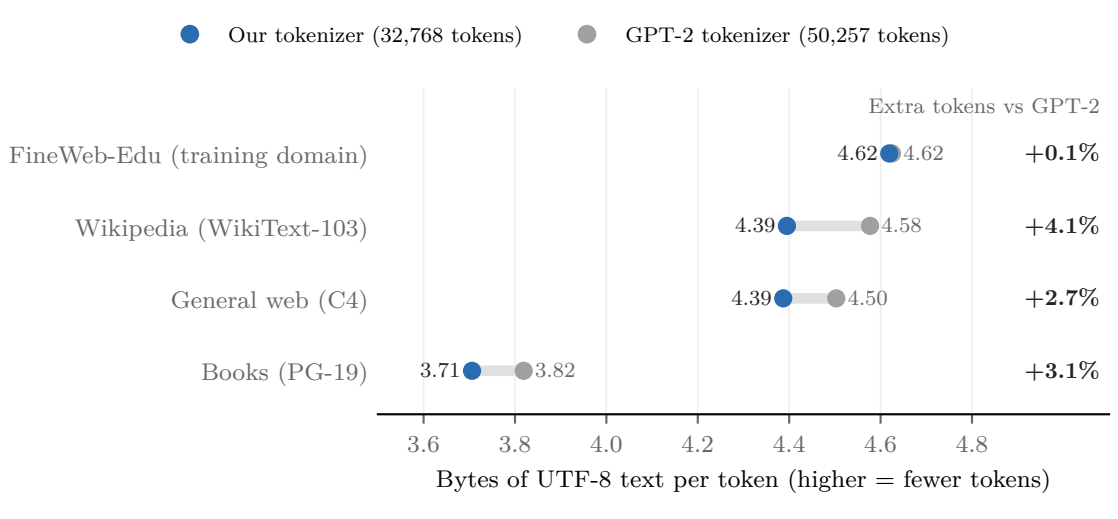
Our tokenizer (32,768 tokens)
GPT-2 tokenizer (50,257 tokens)
Extra tokens vs GPT-2
FineWeb-Edu (training domain)
4.62
4.62
+0.1%
Wikipedia (WikiText-103)
4.39
4.58
+4.1%
General web (C4)
4.39
4.50
+2.7%
Books (PG-19)
3.71
3.82
+3.1%
3.6
3.8
4.0
4.2
4.4
4.6
4.8
Bytes of UTF-8 text per token (higher = fewer tokens)