Decoder-Only Transformer from Scratch

Technical Report

Yash Singh

November 2026

In my earlier projects, like LinkLens-V2 and FACTIFY, I used transformer models through libraries, and I always wanted to understand how they actually work underneath. So instead of using the libraries as they are, I built every part from scratch, from the tokenizer to the evaluation, and trained it at two sizes. This report documents the final models and their results; the blog series describes how I got there, including the mistakes along the way.

Abstract

This report describes two decoder-only transformer language models, with 110M and 336M parameters, built from scratch, including a byte-level BPE tokenizer, rotary position embeddings, multi-head causal attention with a KV cache, and a tested training and evaluation pipeline. Both were trained on FineWeb-Edu (2.5B and 7.0B tokens) on two NVIDIA A30 GPUs, in about 4.6 and 34.7 hours. On held-out FineWeb-Edu text they reach 1.000 and 0.897 bits per byte, better than the size-matched GPT-2 models (1.009 and 0.926), and they outperform them on science and knowledge benchmarks (ARC, SciQ, OpenBookQA), where the 336M also beats GPT-2 large. On other kinds of text, on commonsense benchmarks and on LAMBADA they fall behind, and their 8-task zero-shot average is about one point below GPT-2 at both sizes. The results follow the training data: an educationally filtered corpus buys knowledge and in-domain modelling at the cost of breadth.

Code: github.com/Kevhand/decoder-only-transformer-from-scratch
Models: huggingface.co/Kevhand · **Blog:** yashksingh.hashnode.dev

Contents

1 Summary **1**

2 Tokenizer **1**
2.1 Design 2
2.2 Training 2
2.3 Encoding and integration 2
2.4 Comparison with other tokenizers 3
2.5 Compression on the evaluation text 3
2.6 Why results are reported in bits per byte 4
2.7 What the vocabulary costs the model 4
2.8 Limitations 5

3 Model architecture **5**
3.1 Configuration 5
3.2 Forward pass 5
3.3 Components 6
3.4 Initialization 6
3.5 Parameters 7
3.6 Verification 7

4 Data **7**

5 Training **8**
5.1 Setup 9
5.2 The runs 10

6 Evaluation **11**
6.1 Language modelling: bits per byte 11
6.2 Bits per byte by kind of text 11
6.3 Zero-shot benchmarks 11
6.4 Generation 12
6.5 Generation speed 12

7 Using context **13**
7.1 LAMBADA: the last word of a passage 13
7.2 SQuAD: reading comprehension 13
7.3 Knowing, reading, and following the context 14

8 Engineering problems and fixes **15**

9 Limitations and next steps **17**
9.1 Limitations 17
9.2 Next steps 17

10 Reproducibility **18**
10.1 Code layout 18
10.2 Commands 18
10.3 Artifacts 19

1 Summary

This project builds a GPT-style language model from scratch, tokenizer included, and trains it at two sizes on two GPUs. Apart from PyTorch's tensor operations and its fused attention kernel, every part is our own code: a byte-level BPE tokenizer, token embeddings, rotary position embeddings, multi-head causal attention with a KV cache, LayerNorm, the feed-forward network, and the training and evaluation pipeline, each with its own tests.

	110M	**336M**
Parameters	110.3M	335.9M
Layers × d_model	12 × 768	24 × 1,024
Training data	2.5B tokens of FineWeb-Edu	7.0B tokens of FineWeb-Edu
Training time on 2× A30	≈ 4.6 h at 35% MFU	≈ 34.7 h at 39% MFU
Held-out FineWeb-Edu, bits/byte	**1.000** (GPT-2 small: 1.009)	**0.897** (GPT-2 medium: 0.926; large: 0.880)
8-task zero-shot average	40.6 (GPT-2 small: 41.5)	44.9 (GPT-2 medium: 46.0)

Main results.

- On text like their training data, both models beat the size-matched GPT-2, and the 336M comes within 0.017 bits/byte of GPT-2 large, a model 2.3× its size.
- On science and school-knowledge benchmarks (ARC-Easy, ARC-Challenge, SciQ, OpenBookQA), both models beat the size-matched GPT-2, and the 336M beats GPT-2 large on all four.
- They are weaker on narrative and general text: 0.06–0.11 bits/byte behind on Wikipedia and C4, behind on HellaSwag, PIQA and WinoGrande, and 9–11 points behind on LAMBADA. The overall benchmark average ends about 1 point below GPT-2 at both sizes.
- Given a short, explicit passage, both use it like GPT-2 does (100% with a supporting passage); they extract less from longer reading-comprehension passages and answer in full sentences rather than short spans.

The pattern follows the data: FineWeb-Edu is filtered for educational content, which buys knowledge and in-domain modelling at the cost of breadth. The rest of this report covers the tokenizer, the architecture, the data and training setup, every evaluation, and the engineering problems found along the way.

2 Tokenizer

Both models use a byte-level BPE tokenizer written from scratch [1, 2]: a 32,768-token vocabulary trained on FineWeb-Edu. It compresses held-out FineWeb-Edu to **about 4.6 bytes per token, on par with GPT-2's 50,257-token tokenizer** despite a vocabulary 35% smaller, and encodes any UTF-8 text without unknown tokens.

2.1 Design

Component	Choice
Base vocabulary	The 256 byte values, at IDs 0–255 (token ID = byte value)
Special tokens	`[PAD]` 256, `[BOS]` 257, `[EOS]` 258, `[UNK]` 259. `[UNK]` is never produced, since every byte is in the base vocabulary
Vocabulary size	32,768 (256 bytes + 4 special tokens + 32,508 merges)
Pre-tokenizer	Hand-written and Unicode-aware: splits letters, numbers and other symbols by Unicode category; attaches a leading space to the following word, number or symbol run; keeps newlines as their own chunk; splits the contractions 's, 're, 've, 'll, 'd, n't and 'm from the word they end
Merges	Learned within pre-token chunks only, never across them; applied at encode time lowest rank first
Stored state	Only the ordered merges, as pairs of token IDs; the vocabulary is rebuilt from them on load
Training data	A sample of FineWeb-Edu (sample-10BT)

The pre-tokenizer follows the same idea as GPT-2's regex [2] but differs in one visible way: it splits off "n't" rather than "'t", so don't → "Do" + "n't" and aren't → " are" + "n't", keeping the base word whole.

2.2 Training

- Each document is pre-tokenized separately, in parallel worker processes, so chunks never span documents.
- Training runs on the **distinct** chunks, each weighted by its frequency. This produces the same merges as training on every occurrence, about 6.6× faster.
- Pair counts are updated incrementally: an index from each pair to the chunks containing it means a merge only recounts the affected chunks. The most frequent pair is taken from a max-heap with lazy deletion (stale entries are skipped when popped).
- Chunks are processed in sorted order, so training is deterministic.

2.3 Encoding and integration

- `encode` returns raw token IDs. A chunk-to-IDs cache avoids re-running BPE on repeated chunks, making encoding about 2.4× faster.
- The training data is packed as `[BOS] document [EOS]`, back to back, into 1,024-token sequences, so no padding is used. The model masks the `[PAD]` logit, so it can never be predicted.
- The saved tokenizer records its format version, pre-tokenizer version and contraction list, and loading refuses a mismatch, so the same text can never be split differently from the training data.
- Verified by exact encode–decode round trips, and by re-encoding 300 validation documents and checking they match the token shards exactly.

2.4 Comparison with other tokenizers

The same texts were encoded with each tokenizer. Compression is in bytes of UTF-8 text per token (higher means fewer tokens for the same text); every tokenizer round-tripped all four sets exactly.

Tokenizer	Vocabulary	FineWeb-Edu	Python code	Hindi	Numbers	Vocabulary used on FineWeb-Edu	Encoding speed (MB/s, 1 process)
Ours	32,768	4.58	2.07	2.14	2.68	90.7%	0.89
GPT-2	50,257	4.61	2.25	1.68	2.86	74.1%	14.13
LLaMA-2 [3]	32,000	3.99	3.60	2.42	1.63	57.5%	2.55
GPT-4 (cl100k)	100,277	4.69	4.58	2.62	2.40	40.9%	5.53

Test sets: FineWeb-Edu documents from far past the tokenizer's training sample, Python standard-library source, Hindi Wikipedia articles, and generated number-heavy text (prices, dates, decimals). GPT-2 and GPT-4 were run through tiktoken and LLaMA-2 through Hugging Face transformers.

- **On its training domain, the tokenizer matches GPT-2 with a smaller vocabulary** (4.58 vs 4.61 bytes per token), and it uses 90.7% of its vocabulary on FineWeb-Edu, against 74.1% for GPT-2, so few entries are wasted.
- **It beats GPT-2 on Hindi** (2.14 vs 1.68), but is behind LLaMA-2 and GPT-4, whose training data is far more multilingual.
- **It is weakest on code**: runs of spaces are split into one token per space (see Limitations), which matters most for indentation.
- **Encoding speed** reflects the implementation, not the algorithm: ours is pure Python, while tiktoken is written in Rust.

On a typical English sentence the splits are close to GPT-2's, differing mainly in contractions:

```
"Jane Austen’s novels weren't only about romance; they're about independence."
Ours   18 tokens  'J' 'ane' ' Aust' 'en' '’' 's' ' novels' ' were' "n't" ' only' ' about' ' romance'
';' ' they' "'re" ' about' ' independence' '.'
GPT-2  18 tokens  'Jane' ' Aust' 'en' <e2 80> <99> 's' ' novels' ' weren' "'t" ' only' ' about' '
romance' ';' ' they' "'re" ' about' ' independence' '.'
```

2.5 Compression on the evaluation text

The bits-per-byte evaluations also recorded how many tokens each tokenizer needed for exactly the same text, which gives a second comparison with GPT-2 on the four kinds of text the models were scored on.

The tokenizer shows the same domain pattern as the models: it is as efficient as GPT-2's on the educational web text it was trained on, and less efficient on everything else. Bits per byte is unaffected by token counts, but more tokens means more computation per document and less text per 1,024-token context.

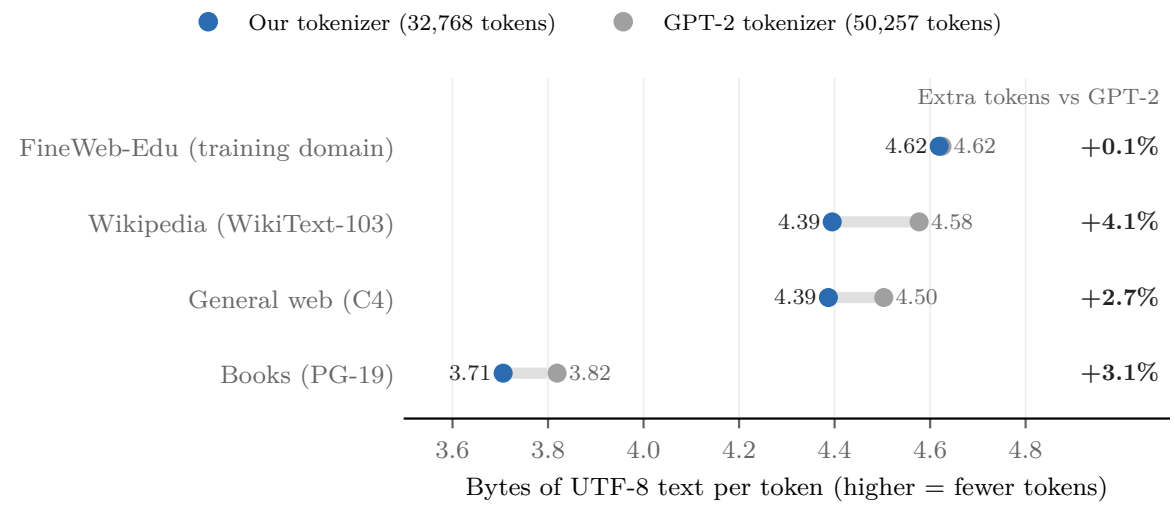


Figure 1: Bytes of UTF-8 text per token on the evaluation text: our tokenizer matches GPT-2 on FineWeb-Edu but needs 2.7–4.1% more tokens on other text. Source: domain_results.json (token counts from the bits-per-byte evaluation; all GPT-2 sizes share one tokenizer, as do both of our models).

2.6 Why results are reported in bits per byte

Loss and perplexity are measured per token, and a token covers a different amount of text in each tokenizer, so they cannot be compared between models with different tokenizers. Bits per byte divides the total loss by the number of UTF-8 bytes in the text, which is identical for every model:

$$\text{bits per byte} = \frac{\text{total loss (nats)}}{\ln 2 \times \text{bytes of text}} = \frac{\text{loss per token}}{\ln 2 \times \text{bytes per token}}$$

The difference matters in these results. On the books text, our 110M model has a lower loss per token than GPT-2 small (3.659 vs 3.732 nats, perplexity 38.8 vs 41.8), which would suggest it predicts books better. But our tokenizer splits that text into 3.1% more tokens, each slightly easier to predict. Per byte, our model is worse: 1.424 vs 1.410 bits per byte. Comparing per-token numbers would have reversed the conclusion.

2.7 What the vocabulary costs the model

Every vocabulary entry needs a row in the token embedding matrix, which is shared with the output layer (weight tying), so it is counted once.

Model	Token embedding matrix	Parameters	Share of the model
Ours (110M)	32,768 × 768	25.2M	22.8%
GPT-2 small (124M)	50,257 × 768	38.6M	31.0%
Ours (336M)	32,768 × 1,024	33.6M	10.0%
GPT-2 medium (355M)	50,257 × 1,024	51.5M	14.5%

At the same compression on FineWeb-Edu, the smaller vocabulary uses 13.4M fewer embedding parameters than GPT-2 small. The share shrinks as models grow, since the embedding grows with the model width while the transformer layers grow with width squared.

2.8 Limitations

- **Whitespace runs:** a space not followed by a letter, number or symbol becomes its own chunk, so indentation costs one token per space. This is the main reason code compresses worse than with GPT-2 (2.07 vs 2.25 bytes per token).
- **Non-English scripts:** the vocabulary is fitted to English educational text. Hindi averages 2.14 bytes per token, under one Devanagari character (3 bytes) per token, so much of it falls back to raw bytes.
- **Numbers:** the pre-tokenizer keeps digit runs together, but BPE then merges digits by frequency, so a number's split depends on its digits (2026 becomes " 20" + "26").
- **Speed:** pure-Python encoding at under 1 MB/s per process; tokenizing the 7B-token training set relied on many parallel worker processes.

3 Model architecture

Both models are built from the same code: a decoder-only transformer [2, 4] with pre-LayerNorm blocks [5, 6], rotary position embeddings on queries and keys [7], a GELU feed-forward network [8], and an output layer tied to the token embedding [9]. Only the configuration differs. Every component is implemented from scratch in PyTorch (no nn.Transformer, nn.LayerNorm or nn.MultiheadAttention); the only fused library call is torch.nn.functional.scaled_dot_product_attention.

3.1 Configuration

	110M	**336M**
Vocabulary	32,768	32,768
d_model	768	1,024
Decoder blocks	12	24
Attention heads × head_dim	12 × 64	16 × 64
FFN hidden size d_ff	3,072 (4×)	4,096 (4×)
Context length	1,024	1,024
Total parameters	110,254,592	335,898,624
Non-embedding parameters	85,088,768	302,344,192
Embedding share of parameters	22.8%	10.0%

3.2 Forward pass

```
token IDs (B, T)
  → token embedding                                        (B, T, d_model)
  → N × decoder block:
        h = x + W_O · Attention(RoPE(Q), RoPE(K), V),  with Q, K, V = projections of LN₁(x)
        y = h + FFN(LN₂(h))
  → final LayerNorm
  → LM head (weight tied to the embedding, own bias)   (B, T, vocab)
  → PAD logit set to −∞
```

3.3 Components

Component	Implementation
Token embedding	nn.Embedding with padding_idx = PAD (ID 256); PAD row initialized to zero
Q, K, V projections	Three separate d_model → d_model linear layers with bias, reshaped to (B, heads, T, head_dim)
Positional information	RoPE on Q and K only, base 10,000, interleaved (even, odd) pairs, 32 pairs per head; cos/sin tables precomputed for 1,024 positions in float32, registered as non-persistent buffers (not saved in checkpoints)
Attention	Multi-head causal attention via scaled_dot_product_attention, scale $1/\sqrt{64}$; a manual path (explicit scores, $-\infty$ mask, softmax) with identical results is kept for inspecting attention weights
Output projection	W_O, d_model → d_model with bias
Feed-forward	Linear (d_model → 4·d_model) → exact GELU → Linear (4·d_model → d_model), both with bias
Normalization	Custom LayerNorm (population variance, ε = 1e-5, learnable γ and β): pre-LN before attention and before the FFN in every block, plus one final LayerNorm after the last block
Output layer	nn.Linear sharing its weight matrix with the token embedding (saves 25.2M / 33.6M parameters), with its own bias; the PAD logit is masked to $-\infty$ so PAD is never predicted and its row receives no gradient through the head

KV cache. Each block caches its keys after RoPE and its values, so a generation step only projects, rotates and attends for the new tokens. RoPE is applied to the new Q and K at their absolute offset (past_len), and attention builds its causal mask from absolute positions, so a rectangular (new tokens × all keys) mask is supported. Attention has three paths: is_causal=True for a fresh pass, no mask for a single new token (it may attend to every cached key), and an explicit position mask for a multi-token chunk on top of a cache. Mismatched cache lengths and cache lists of the wrong length raise errors.

Mixed precision. Training and bf16 evaluation run under torch.autocast with bfloat16. RoPE's tables are float32 and are not cast down, so the rotation itself runs in float32 and returns float32 Q and K; attention then runs in bf16 under autocast. Measured on head_dim 64 across all 1,024 positions, this keeps the drift of a fixed-gap Q·K score at 0.12, against 0.19 when the tables are cast to bf16 first (0.0002 in pure float32).

3.4 Initialization

GPT-2 style:

Parameter	Initialization
Token embedding (and so the tied LM head weight)	Normal(0, 0.02), PAD row zero
Linear weights	Normal(0, 0.02)
Residual projections (W_O and the FFN's second layer)	Normal(0, 0.02 / $\sqrt{(2N)}$): 0.0041 for the 110M (N = 12), 0.0029 for the 336M (N = 24)
Linear biases	Zero
LayerNorm γ / β	One / zero

The embedding is initialized last, after the linear layers, because the tied LM head is an nn.Linear

and would otherwise overwrite the shared matrix and its zero PAD row. With this initialization the initial logits have a standard deviation of about 0.16 and the initial loss matches the uniform baseline $\ln(V - 1)$; with PyTorch's default N(0, 1) embedding the logits' standard deviation was about 8 and the initial loss on a 1,000-token test model was 56.4 instead of 6.9.

3.5 Parameters

Part	110M	336M
Token embedding / tied LM head weight	25.17M (22.8%)	33.55M (10.0%)
Attention (Q, K, V, W_O), all blocks	28.35M (25.7%)	100.76M (30.0%)
Feed-forward, all blocks	56.67M (51.4%)	201.45M (60.0%)
LayerNorms (2 per block + final)	0.04M	0.10M
LM head bias	0.03M	0.03M
Per decoder block	7.09M	12.60M

3.6 Verification

Each component has its own test file, and the whole suite passes on CPU in about 93 seconds. The tests check behaviour, not just shapes, and each property test has a negative control so it can fail:

- **Against references:** LayerNorm matches nn.LayerNorm in the forward pass and in all gradients (with initial and with random γ/β); the FFN matches a hand computation; the fused attention path matches the manual path in outputs (difference 2.4e-7) and gradients.
- **RoPE:** each pair is rotated by position × frequency (hand computation); vector length is preserved; position 0 is the identity; Q·K depends only on the relative position (same gap: difference 6e-8, different gap: 0.149); out-of-range, negative-offset and odd-head_dim inputs are rejected.
- **Causality:** changing token 8 leaves the outputs and logits at positions 0–7 unchanged and changes every position from 8 on, in both attention paths.
- **KV cache:** token-by-token and uneven chunked processing (3 + 4 + 5 tokens) match a full pass at block and model level, in both attention paths.
- **Model:** Pre-LN wiring (with sublayer outputs zeroed, a block is exactly the identity); weight tying survives a save/load round trip; the PAD row stays exactly zero through AdamW steps with weight decay; initial loss within 0.1 of $\ln(V - 1)$; a one-batch overfit goes from 6.89 to 0.0096 in 150 steps.

On the trained checkpoints, cached and uncached generation pick the same tokens over 64 steps, with a maximum logit difference of 2.2e-5 (110M) and 2.3e-5 (336M) in float32.

4 Data

Both models are trained on **FineWeb-Edu** [10] (HuggingFaceFW/fineweb-edu, the sample-10BT subset): English web pages filtered by a classifier for educational value, streamed from the Hugging Face Hub.

	110M	**336M**
Training tokens	2,499,805,184 (2.50B)	6,999,769,088 (7.00B)
Passes over the data	1	1
Tokens per parameter	22.7	20.8

Both models get roughly 20 tokens per parameter, about the compute-optimal ratio from the Chinchilla scaling work [11], and every training token is seen once.

Tokenization and packing. The documents are tokenized with the project's byte-level BPE tokenizer (see Tokenizer), in parallel worker processes. Each document becomes [BOS] + tokens + [EOS], and documents are packed back to back with no padding, so a 1,024-token training window can span several documents and no compute is spent on PAD. Tokens are stored as uint16 (the 32,768-token vocabulary fits) in .npy shards of 100M tokens (200 MB each), with a manifest recording the shards, the token counts, the tokenizer and how many source documents were consumed.

Two datasets, one stream. The 2.5B-token dataset was tokenized first, for the 110M. For the 336M it was extended to 7B tokens by resuming from the manifest: tokenization continued from the first document not yet used, appending new training shards. So the 336M's data is a superset of the 110M's, both start from the same documents, and the validation shard is identical for both runs.

Validation split. The first 10,121,940 tokens of the stream (about 10.1M) form the validation shard, and everything after goes to training shards, so no validation document is ever trained on. During training, validation loss and bits per byte are computed every 250 steps on a fixed slice from the start of this shard: 64 micro-batches, which is 1.05M tokens for the 110M (micro-batch 16) and 0.52M tokens for the 336M (micro-batch 8).

Held-out evaluation text. The final evaluations use separate text: FineWeb-Edu documents from **after document 7,500,000** of the same stream, beyond everything used for training (and for the tokenizer). 1,036 documents (5.0 MB) are scored, each starting with the model's start token and split into non-overlapping 1,024-token windows. This differs from the in-training validation (packed windows that cross document boundaries), so the two sets of numbers are close but not identical: 1.0240 vs 1.0004 bits per byte for the 110M, and 0.9120 vs 0.8966 for the 336M.

Data loading. The shards are memory-mapped and read as one long stream in a shuffled shard order (seeded per epoch). Micro-step k on GPU r reads chunk $k \times \text{world_size} + r$, a pure function of the step, so the two GPUs never read the same tokens and a resumed run continues exactly where it stopped. Each chunk is $\text{micro_batch} \times 1{,}024 + 1$ tokens, giving inputs and next-token targets offset by one.

5 Training

Both models were trained with the same script (train.py) on one server with **2× NVIDIA A30 (24 GB)**, using PyTorch DistributedDataParallel launched by torchrun, bfloat16 autocast, and torch.compile.

5.1 Setup

	110M	**336M**
Tokens per optimizer step	524,288 (2^{19})	524,288 (2^{19})
Per step	16 seq × 1,024 × 2 GPUs × 16 accumulation	8 seq × 1,024 × 2 GPUs × 32 accumulation
Optimizer steps	4,768	13,351
Peak learning rate	6e-4	3e-4
Warmup	143 steps (3%)	400 steps (3%)
Final learning rate	6e-5 (10% of peak)	3e-5 (10% of peak)
Mean throughput	150,796 tokens/s	56,071 tokens/s
Time per step	3.47 s	9.34 s
MFU	35.4%	39.4%
Training time (steps only)	≈ 4.6 h	≈ 34.7 h
Training compute	≈ 1.9×10^{18} FLOPs	≈ 1.6×10^{19} FLOPs

Shared settings:

- **Optimizer:** AdamW with [12] $\beta = (0.9, 0.95)$, $\varepsilon =$ 1e-8 (fused CUDA implementation). Weight decay 0.1 on every matrix (all linear weights and the tied embedding); none on biases or LayerNorm parameters.
- **Schedule:** linear warmup over the first 3% of steps, then cosine decay to 10% of the peak learning rate at the last step.
- **Gradient clipping** at a global norm of 1.0.
- **Loss:** cross-entropy over all positions of the packed windows.
- **Precision:** bf16 autocast for the forward and backward pass; float32 master weights and optimizer state.
- **Checkpoints** every 250 steps for the 110M and every 100 steps for the 336M (the longer run), with resumable data order; weights-only copies kept every 1,000 steps.
- **Monitoring:** training loss, learning rate, gradient norm and throughput every 10 steps; validation loss and bits per byte every 250 steps; a sampled completion of "The most important thing about learning is" every 500 steps.

The 336M uses half the peak learning rate of the 110M, a common adjustment for a larger model, and a micro-batch of 8 so that its activations fit in 24 GB, with twice the gradient accumulation to keep the same 2^{19}-token batch. It was launched with PYTORCH_CUDA_ALLOC_CONF=expandable_segments:True to reduce memory fragmentation.

MFU (model FLOPs utilization [13]) is measured against the A30's 165 TFLOPS bf16 peak per GPU, counting 6N FLOPs per token for the weights plus 12 · layers · d_model · context for attention. Before the 110M run, a 30-step benchmark showed what torch.compile is worth:

110M, 30-step benchmark	**Throughput**	**MFU**	**Peak memory per GPU**
Eager	91,544 tokens/s	21.5%	18.9 GB
torch.compile	153,375 tokens/s	36.0%	15.5 GB

Compiling made training 1.68× faster and used less memory, so both runs used it; uncompiled, the 110M run would have taken about 7.6 hours instead of 4.6.

5.2 The runs

Both runs went from start to finish without a restart, a loss spike or a divergence. The initial loss was 10.55 (110M) and 10.65 (336M), close to ln(32,767) = 10.40 for a uniform guess over every token except PAD. After warmup the gradient norm settled quickly: its median over the run was 0.27 for the 110M and 0.23 for the 336M, and it went above the clipping threshold of 1.0 only a handful of times, all within a few hundred steps of the end of warmup (2 logged steps for the 110M, 4 for the 336M, at most 1.07).

Validation (in-training set)	**110M**	**336M**
After 0.5B tokens (step 1,000)	1.197 bits/byte	1.222 bits/byte
After 1.0B tokens (step 2,000)	1.102	1.084
After 2.5B tokens (step 4,750)	1.025	0.989
Final	**1.024** (loss 3.262)	**0.912** (loss 2.882)

The 336M starts slower, with a lower peak learning rate and a longer warmup, overtakes the 110M at about 0.75B tokens (step 1,500), and at 2.5B tokens, where the 110M's training ended, it is already 0.035 bits/byte better, even though its learning rate at that point is still about three-quarters of its peak (2.3e-4) while the 110M's has fully decayed. The remaining 4.5B tokens take it from 0.989 to 0.912. Both curves are still decreasing at the end: neither model has reached the limit of what more data would give it.

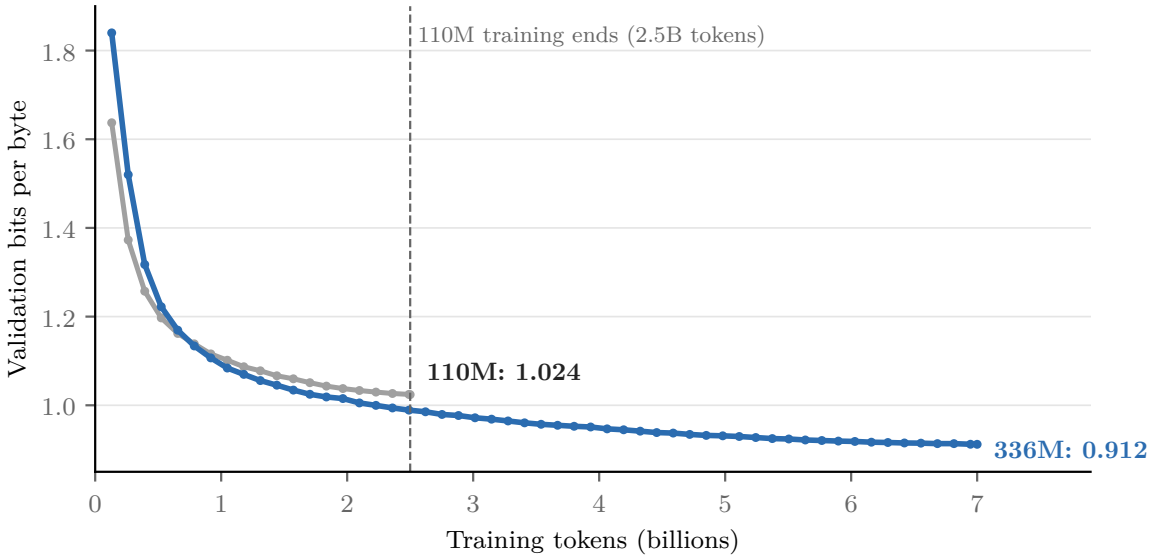


Figure 2: Validation bits per byte during training (in-training validation set, every 250 steps; lower is better). The 336M reaches 0.912; the 110M stops at 1.024 after 2.5B tokens. Source: log_110m.jsonl and log_336m.jsonl, 20 and 54 points.

The completion of the fixed prompt at the end of each run (temperature 0.8, top-k 50):

- **110M:** "The most important thing about learning is that you can't do anything without learning. Learning is all about memorizing and remembering…"
- **336M:** "The most important thing about learning is that it is not easy. You have to change every single aspect of your life. You have to learn t…"

6 Evaluation

Every evaluation runs our two models and the three GPT-2 models from Hugging Face [2] (small 124M, medium 355M, large 774M) on identical data with identical scoring. GPT-2 small is the size match for the 110M and GPT-2 medium for the 336M. GPT-2 was trained by OpenAI on WebText, a broader web corpus than FineWeb-Edu, so these comparisons measure both the models and their training data.

6.1 Language modelling: bits per byte

On the held-out FineWeb-Edu text (1,036 documents, 5.0 MB, from after everything used for training):

Model	Bits/byte	Loss/token	Perplexity	Bytes/token
Ours (110M)	**1.0004**	3.2032	24.61	4.62
GPT-2 small (124M)	1.0093	3.2358	25.43	4.63
Ours (336M)	**0.8966**	2.8708	17.65	4.62
GPT-2 medium (355M)	0.9262	2.9694	19.48	4.63
GPT-2 large (774M)	0.8801	2.8215	16.80	4.63

Bits per byte is the comparable measure, since each model uses its own tokenizer (see Tokenizer for why loss per token can mislead). On text like its training data, the 110M is slightly better than GPT-2 small, and the 336M is 0.030 bits/byte better than GPT-2 medium and within 0.017 of GPT-2 large, a model 2.3× its size.

6.2 Bits per byte by kind of text

The same measurement on three other kinds of text: Wikipedia [14] (WikiText-103 test, 1.2 MB), books [15] (PG-19 test, 5.1 MB, a slice of each book starting 20% in) and general web text [16] (C4 validation, 5.0 MB).

Model	FineWeb-Edu	Wikipedia	Books	General web (C4)
Ours (110M)	**1.0004**	1.1257	1.4244	1.1832
GPT-2 small (124M)	1.0093	**1.0155**	**1.4098**	**1.0968**
Ours (336M)	**0.8966**	0.9980	1.2961	1.0741
GPT-2 medium (355M)	0.9262	**0.9212**	**1.2489**	**1.0138**
GPT-2 large (774M)	0.8801	0.8706	1.2159	0.9681

Outside the training domain the order reverses: GPT-2 is better on all three, by 0.11 (110M vs small) and 0.08 (336M vs medium) bits/byte on Wikipedia, 0.09 and 0.06 on C4, and 0.01 and 0.05 on books. FineWeb-Edu is filtered towards educational web pages, so text that looks different (encyclopedic formatting, fiction, everyday web pages) is less familiar to our models. Scale still helps everywhere: the 336M improves on the 110M by 0.10–0.13 bits/byte in every domain.

6.3 Zero-shot benchmarks

Eight standard tasks from EleutherAI's lm-evaluation-harness [17], run through a wrapper around our model class, in float32, zero-shot. Before trusting the wrapper, it was run on

GPT-2 small next to the harness's own Hugging Face model: all 14 task/metric pairs agreed within 0.1 points.

Model	HellaSwag [18]	PIQA [19]	ARC-e [20]	ARC-c [20]	WinoGrande [21]	SciQ [22]	OpenBookQA [23]	LAMBADA [24]	Average
Ours (110M)	29.5	60.1	**43.1**	**24.1**	51.7	**65.4**	**28.0**	22.8	40.6
GPT-2 small (124M)	**31.1**	**62.6**	39.5	22.7	**51.8**	64.4	27.2	**32.6**	**41.5**
Ours (336M)	35.4	63.3	**48.3**	**28.0**	51.6	**70.4**	**31.8**	30.7	44.9
GPT-2 medium (355M)	**39.4**	**66.4**	43.6	25.1	**53.2**	67.2	30.2	**43.1**	**46.0**
GPT-2 large (774M)	45.3	69.2	46.6	25.1	55.5	69.3	31.2	47.7	48.7

Accuracy in %, acc_norm where the task has it, otherwise acc. Standard errors: about ±0.5 (HellaSwag), ±1.0–1.5 (PIQA, ARC, WinoGrande, SciQ), ±2.0 (OpenBookQA) and ±0.6–0.7 (LAMBADA). Chance is 25% for the four-choice tasks, 50% for PIQA and WinoGrande, 0% for LAMBADA.

The results split cleanly by kind of task. On the **science and school-knowledge** tasks (ARC-Easy, ARC-Challenge, SciQ, OpenBookQA), both of our models beat their size-matched GPT-2, and the 336M beats GPT-2 large on all four (ARC-Easy 48.3 vs 46.6, ARC-Challenge 28.0 vs 25.1, SciQ 70.4 vs 69.3, OpenBookQA 31.8 vs 31.2). That's what FineWeb-Edu's filtering selects for. On **everyday commonsense and narrative** tasks (HellaSwag, PIQA, WinoGrande) GPT-2 is ahead, and on **LAMBADA**, which needs the last word of a story passage, it is far ahead (10 to 12 points). The averages end up about 1 point below GPT-2 at both sizes: 40.6 vs 41.5 and 44.9 vs 46.0.

HellaSwag was also measured with the project's own evaluate.py, which normalizes by tokens instead of bytes: 27.25% (110M), 32.87% (336M), 29.52% (GPT-2 small, matching its published figure), 37.67% (medium) and 43.12% (large). The ordering is the same as in the harness.

6.4 Generation

The same prompts were completed by every model with greedy decoding, temperature sampling (0.8, top-k 50) and nucleus sampling (0.8, top-p 0.9); the full outputs are in results/eval_report.md. Greedy decoding loops for every model at these sizes. Factual prompts show the domain difference: for "The capital of France is", the 336M's greedy completion is "Paris.", the only model to answer directly (GPT-2 medium continues "a city of the future", GPT-2 large "a city of contrasts"), while the 110M produces a circular "the capital of the French Republic".

6.5 Generation speed

Batch 1, tokens per second while generating from a short prompt (A30, bf16; evaluate.py's speed task):

Model	With KV cache	Without cache	Speedup
Ours (110M)	98.4	97.2	1.0×
GPT-2 small (124M)	141.9	117.5	1.2×
Ours (336M)	49.7	50.0	1.0×
GPT-2 medium (355M)	77.1	62.5	1.2×
GPT-2 large (774M)	53.6	37.5	1.4×

At batch size 1 the time per token is dominated by launching many small GPU operations rather than by arithmetic. Hugging Face's GPT-2 launches fewer of them per layer (one fused Q/K/V projection, for example, where ours has three), so it is faster at the same size, and the KV cache saves comparatively more of its time. The KV cache correctness and the fix that made our cache faster than recomputing are under Engineering problems and fixes.

7 Using context

The benchmarks above mostly test what a model already knows. Three further tests (context_eval.py) check whether the models **use** information given in the prompt. The same data and scoring are used for all five models.

7.1 LAMBADA: the last word of a passage

5,153 passages (OpenAI version) [24] whose last word is easy to guess from the whole passage but hard from the last sentence alone. A passage counts as correct only if greedy decoding produces every token of the word.

Model	Accuracy	Perplexity of the target word
Ours (110M)	22.38%	158.9
GPT-2 small (124M)	**30.93%**	**38.9**
Ours (336M)	30.55%	51.7
GPT-2 medium (355M)	**41.41%**	**17.8**
GPT-2 large (774M)	47.31%	13.0

This is our models' weakest result: about 9 and 11 points behind the size-matched GPT-2, and the 336M only reaches GPT-2 small. LAMBADA passages are taken from novels, and keeping track of characters and events through a story is a skill FineWeb-Edu's educational pages rarely train; the books bits-per-byte result points the same way.

7.2 SQuAD: reading comprehension

1,000 SQuAD v1.1 questions [25], each with its Wikipedia paragraph and two worked examples (question → short answer) in front. The model generates an answer until the first newline or 20 tokens. Each question is also asked closed book, without the paragraph, so the gain from reading is visible.

Model	EM, with passage	F1, with passage	F1, without passage	F1 gain from passage
Ours (110M)	0.1%	8.3%	4.6%	+3.6 pts
GPT-2 small (124M)	9.3%	19.8%	4.8%	+15.0 pts
Ours (336M)	3.2%	14.6%	6.3%	+8.3 pts
GPT-2 medium (355M)	24.0%	36.4%	8.0%	+28.4 pts
GPT-2 large (774M)	32.3%	45.7%	8.2%	+37.5 pts

Exact match and F1 punish one thing our models do: they answer in full sentences. Their answers average 13.4 (110M) and 11.8 (336M) words against 4–5 for GPT-2, so a correct answer

inside a sentence still scores near zero. A looser check, whether a correct answer appears anywhere in the output as whole words, separates **not finding** the answer from **phrasing** it badly:

Model	Answer found, with passage	Without passage	Gain from passage	Answer length (words)
Ours (110M)	8.9%	0.8%	+8.1 pts	13.4
GPT-2 small (124M)	16.4%	1.2%	+15.2 pts	5.0
Ours (336M)	**18.3%**	2.0%	+16.3 pts	11.8
GPT-2 medium (355M)	31.1%	3.2%	+27.9 pts	4.0
GPT-2 large (774M)	42.5%	3.0%	+39.5 pts	4.2

With the looser check the 336M finds the answer more often than GPT-2 small (18.3% vs 16.4%), but it still finds it far less often than GPT-2 medium (31.1%). Part of the EM gap is format, then, and most of it is real: our models also extract less from the passage. The two worked examples did not get either model to copy the short-answer format; its training text contains few question-and-answer pairs of that kind. A typical 336M answer: *"Super Bowl 50 was held on February 7, 2016, at the San Francisco Bay Area. The game…"* (correct, and scored 0 by exact match).

7.3 Knowing, reading, and following the context

25 simple fact questions ("What is the capital of Japan?": Tokyo or Osaka), each asked three ways. The model is scored by which of the two answers it finds more likely: with no passage (does it know?), after a passage stating the true answer (does it read?), and after a passage stating the false one (does it follow the context or its memory?).

Model	Knows it (no passage)	Reads it (true passage)	Follows a false passage
Ours (110M)	80%	100%	92%
GPT-2 small (124M)	68%	96%	96%
Ours (336M)	80%	100%	96%
GPT-2 medium (355M)	80%	100%	100%
GPT-2 large (774M)	84%	100%	100%

With 25 questions, one question is 4 points, so these are coarse. The 110M knows more of these facts than GPT-2 small (80% vs 68%) and as many as GPT-2 medium, consistent with its strength on the knowledge benchmarks. All models read a supporting passage almost perfectly. Given a passage that contradicts what they know, all of them mostly follow the passage; the 110M keeps its own answer on 2 of the 25 questions, the 336M and GPT-2 small on 1, GPT-2 medium and large on none.

Taken together: with a short, explicit passage our models use the context as well as GPT-2 does, but on longer and less structured context (stories, a Wikipedia paragraph answered in free form) they are clearly weaker. That's consistent with their training data: lots of explanatory text, few narratives, and little question-and-answer formatting.

8 Engineering problems and fixes

The problems that changed the code or the results, in the order they came up.

Problem	How it showed	Fix	Effect
Initial loss far too high	Untrained 1,000-token test model started at a loss of 56.4 instead of $\ln(1000) = 6.9$	GPT-2-style initialization (std 0.02, scaled residual projections), embedding initialized last because of weight tying	Initial loss 6.90, logits' std $8.2 \to 0.16$
PAD row of the tied matrix could change	padding_idx protects the embedding lookup, but the tied LM head still gave the PAD row gradients	PAD logit set to $-\infty$, PAD excluded from targets, row initialized to zero	Row stays exactly 0 through AdamW with weight decay (tested)
A test passed while printing NaN	Cache-equivalence max difference printed nan: $-\infty - (-\infty)$ on the PAD column, while allclose treats equal infinities as close	Compare non-PAD logits only	Real differences again (2.8e-7)
KV cache could not handle generation	First RoPE always started at position 0; the attention mask was a sliced square, wrong when Q is shorter than K	Absolute offsets in RoPE, keys cached after rotation, mask built from absolute positions	Token-by-token and chunked runs match a full pass to ~1e-7
Attention memory	Manual attention stores (B, heads, T, T) weights: 48 MB (110M) / 64 MB (336M) per layer per sequence in float32	scaled_dot_product_ attention, with the manual path kept for inspection	Same results (2.4e-7), far less memory
Slow eager training	110M benchmark at 91.5k tokens/s, MFU 21.5%	torch.compile	153k tokens/s, MFU 36.0%, 1.68× faster, less memory
336M memory	24 GB per A30	Micro-batch 8 with 32 accumulation steps; expandable_segments allocator	Same 2^{19}-token batch, 39.4% MFU
KV cache slower than no cache	Full-context generation: 89.4 vs 92.9 tokens/s (110M), 46.3 vs 47.9 (336M)	A single new token skips the mask (it may see every cached key), instead of building aranges and a mask in every layer at every step	96.8 vs 93.1 (1.04×) and 49.8 vs 49.0 (1.02×); up to 1.08× in the last quarter of the context
bf16 vs float32 benchmark scores	Harness results in bf16 differed from float32 by up to 1.8 points (LAMBADA)	All reported harness numbers re-run in float32; our wrapper checked against the harness's own GPT-2	Wrapper within 0.1 points on all 14 task/ metric pairs
SQuAD exact match undercounted our models	Full-sentence answers containing the right answer scored 0	Added an "answer found" check and answer length	336M: 3.2% EM, but 18.3% answer found

The tokenizer had its own set of speed problems (weighted distinct-chunk training, about 6.6×

faster; an encode cache, about 2.4×; a lazy-deletion heap and incremental pair counts), described under Tokenizer.

One decision that wasn't a bug: RoPE keeps its rotation in float32 under bf16 autocast instead of casting its tables down. Measured on head_dim 64 across 1,024 positions, this lowers the drift of a fixed-gap Q·K score from 0.19 to 0.12.

The code fixes are covered by tests in the final suite (the KV-cache speed fix by the benchmark in results/kv_cache_report.md), and the suite was run before and after the attention change.

9 Limitations and next steps

9.1 Limitations

- **Narrow training data.** FineWeb-Edu makes both models strong on educational and science text and on knowledge benchmarks, but weaker than GPT-2 on fiction, Wikipedia-style text and general web text (0.06–0.11 bits/byte), on commonsense tasks, and especially on LAMBADA (9–11 points behind the size-matched GPT-2).
- **Base models only.** Neither model is instruction-tuned. They don't follow a question-and-answer format from two examples (SQuAD answers average 12–13 words), and greedy decoding falls into repetition, as it does for GPT-2 at these sizes.
- **Context length.** 1,024 tokens. The RoPE tables stop at position 1,023 and longer inputs are rejected; no context extension was tried.
- **Tokenizer.** Each run of spaces costs a token per space, Hindi and other non-Latin scripts fall back to bytes, and numbers split irregularly (see Tokenizer).
- **One run per size.** No repeated seeds and no ablations, so small differences (a point on a benchmark, a few questions in the 25-question test) are within noise.
- **Inference speed.** At batch 1, our implementation is slower than Hugging Face's GPT-2 at the same size (98 vs 142 tokens/s for the 110M against GPT-2 small), and the KV cache gives only 1.02–1.08×, because time goes to launching many small operations rather than to computation.
- **The comparison confounds model and data.** GPT-2 differs in its training data, its tokenizer and its training length, so the comparisons show where these models stand, not which architectural choice causes the difference.

9.2 Next steps

- **Inference:** fuse the Q/K/V projections into one matrix multiplication, and capture the decoding step in a CUDA graph to remove the per-operation launch overhead.
- **Ablations at the 110M scale:** SwiGLU instead of the GELU FFN at equal parameters, learned or sinusoidal positions instead of RoPE, untied output weights, and the residual-projection initialization.
- **Data:** a mix that adds narrative and general web text to FineWeb-Edu, to test whether it closes the LAMBADA and off-domain gaps without losing the knowledge-benchmark lead.
- **Interpretability:** weights-only checkpoints were kept every 1,000 steps for both runs, and the manual attention path returns attention weights, so how heads and features develop during training can be studied directly.

- **Instruction tuning** of the 336M on a small question-and-answer dataset, and a longer context.

10 Reproducibility

10.1 Code layout

Path	Contents
Tokenizer/	byte_level_bpe.py (the tokenizer), train_tokenizer.py, compare_tokenizers.py, tokenizer.json
Embedding/	TokenEmbedding, QKV_Projection
Positional_Encoding/	ROPE
Attention_Mechanism/	MultiHeadCausalAttention
FeedForward_Network/	FFN
LayerNorm/	LayerNorm
Decoder_Blocks/	DecoderBlock, TransformerModel
tests/	one test file per component; run_all.py runs the suite
Tokenizer/tokenize_dataset.py	streams FineWeb-Edu and writes the uint16 shards and manifest
train.py, run_336m.sh	DDP training, checkpoints, logging, benchmark mode
evaluate.py, domain_eval.py, harness_eval.py, context_eval.py	the evaluations against GPT-2
check_kv_cache.py	KV-cache correctness and generation speed
plot_results.py, export_artifacts.py	charts and the exported final models
results/	every evaluation report (JSON and Markdown) and chart
logs/	training logs, console output and run summaries of both models
docs/	this report as a PDF, and its LaTeX source

10.2 Commands

```
# Data: 2.5B tokens, later extended to 7B with --resume
python Tokenizer/tokenize_dataset.py --tokenizer Tokenizer/tokenizer.json --out-dir data
--max-tokens 2.5e9
python Tokenizer/tokenize_dataset.py --tokenizer Tokenizer/tokenizer.json --out-dir data
--max-tokens 7e9 --resume

# 110M (the script's defaults)
torchrun --standalone --nproc_per_node=2 train.py --compile

# 336M
PYTORCH_CUDA_ALLOC_CONF=expandable_segments:True torchrun --standalone --nproc_per_node=2 train.py \
  --n-layers 24 --d-model 1024 --n-heads 16 --d-ff 4096 --lr 3e-4 --max-tokens 7e9 \
  --micro-batch 8 --compile --out-dir checkpoints_350m --ckpt-every 100

# Tests and evaluations
python -m tests.run_all
python evaluate.py
python domain_eval.py
```

```
python harness_eval.py
python context_eval.py
python check_kv_cache.py --label after
python plot_results.py
```

10.3 Artifacts

- Training logs: logs/log_110m.jsonl and logs/log_336m.jsonl, with the training loss, learning rate, gradient norm and throughput every 10 steps and the validation results every 250 steps; the full console output of both runs (logs/train_stdout_110m.txt, logs/train_stdout_336m.txt).
- Run summaries: logs/summary_110m.json and logs/summary_336m.json (configuration, token counts, validation curve).
- Evaluation results: results/ (eval, domain, harness, context and KV-cache reports as JSON and Markdown, plus the charts).
- The checkpoints of both models are backed up on the Hugging Face Hub (Kevhand/decoder-110m, and Kevhand/decoder-350m for the 336M, 22.8 GB), and the tokenized training shards at Kevhand/fineweb-edu-32k-tokenized.

Random seed 1337. Training data order is a deterministic function of the step, so a run resumed from a checkpoint continues on exactly the same data.

References

[1] Rico Sennrich, Barry Haddow, and Alexandra Birch. Neural machine translation of rare words with subword units. In *Proceedings of the 54th Annual Meeting of the Association for Computational Linguistics*, 2016.

[2] Alec Radford, Jeffrey Wu, Rewon Child, David Luan, Dario Amodei, and Ilya Sutskever. Language models are unsupervised multitask learners. Technical report, OpenAI, 2019.

[3] Hugo Touvron, Louis Martin, Kevin Stone, et al. Llama 2: Open foundation and fine-tuned chat models. *arXiv preprint arXiv:2307.09288*, 2023.

[4] Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N. Gomez, Łukasz Kaiser, and Illia Polosukhin. Attention is all you need. In *Advances in Neural Information Processing Systems*, 2017.

[5] Jimmy Lei Ba, Jamie Ryan Kiros, and Geoffrey E. Hinton. Layer normalization. *arXiv preprint arXiv:1607.06450*, 2016.

[6] Ruibin Xiong, Yunchang Yang, Di He, Kai Zheng, Shuxin Zheng, Chen Xing, Huishuai Zhang, Yanyan Lan, Liwei Wang, and Tie-Yan Liu. On layer normalization in the transformer architecture. In *Proceedings of the 37th International Conference on Machine Learning*, 2020.

[7] Jianlin Su, Yu Lu, Shengfeng Pan, Ahmed Murtadha, Bo Wen, and Yunfeng Liu. RoFormer: Enhanced transformer with rotary position embedding. *arXiv preprint arXiv:2104.09864*, 2021.

[8] Dan Hendrycks and Kevin Gimpel. Gaussian error linear units (GELUs). *arXiv preprint arXiv:1606.08415*, 2016.

[9] Ofir Press and Lior Wolf. Using the output embedding to improve language models. In *Proceedings of the 15th Conference of the European Chapter of the Association for Computational Linguistics*, 2017.

[10] Guilherme Penedo, Hynek Kydlíček, Loubna Ben Allal, Anton Lozhkov, Margaret Mitchell, Colin Raffel, Leandro Von Werra, and Thomas Wolf. The FineWeb datasets: Decanting the web for the finest text data at scale. In *Advances in Neural Information Processing Systems, Datasets and Benchmarks Track*, 2024.

[11] Jordan Hoffmann, Sebastian Borgeaud, Arthur Mensch, et al. Training compute-optimal large language models. In *Advances in Neural Information Processing Systems*, 2022.

[12] Ilya Loshchilov and Frank Hutter. Decoupled weight decay regularization. In *International Conference on Learning Representations*, 2019.

[13] Aakanksha Chowdhery, Sharan Narang, Jacob Devlin, et al. PaLM: Scaling language modeling with pathways. *arXiv preprint arXiv:2204.02311*, 2022.

[14] Stephen Merity, Caiming Xiong, James Bradbury, and Richard Socher. Pointer sentinel mixture models. In *International Conference on Learning Representations*, 2017.

[15] Jack W. Rae, Anna Potapenko, Siddhant M. Jayakumar, and Timothy P. Lillicrap. Compressive transformers for long-range sequence modelling. In *International Conference on Learning Representations*, 2020.

[16] Colin Raffel, Noam Shazeer, Adam Roberts, Katherine Lee, Sharan Narang, Michael Matena, Yanqi Zhou, Wei Li, and Peter J. Liu. Exploring the limits of transfer learning with a unified text-to-text transformer. *Journal of Machine Learning Research*, 21 (140):1–67, 2020.

[17] Leo Gao, Jonathan Tow, Baber Abbasi, et al. A framework for few-shot language model evaluation. EleutherAI, lm-evaluation-harness, `https://github.com/EleutherAI/lm-evaluation-harness`, 2023.

[18] Rowan Zellers, Ari Holtzman, Yonatan Bisk, Ali Farhadi, and Yejin Choi. HellaSwag: Can a machine really finish your sentence? In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics*, 2019.

[19] Yonatan Bisk, Rowan Zellers, Ronan Le Bras, Jianfeng Gao, and Yejin Choi. PIQA: Reasoning about physical commonsense in natural language. In *Proceedings of the AAAI Conference on Artificial Intelligence*, 2020.

[20] Peter Clark, Isaac Cowhey, Oren Etzioni, Tushar Khot, Ashish Sabharwal, Carissa Schoenick, and Oyvind Tafjord. Think you have solved question answering? try ARC, the AI2 reasoning challenge. *arXiv preprint arXiv:1803.05457*, 2018.

[21] Keisuke Sakaguchi, Ronan Le Bras, Chandra Bhagavatula, and Yejin Choi. WinoGrande: An adversarial winograd schema challenge at scale. In *Proceedings of the AAAI Conference on Artificial Intelligence*, 2020.

[22] Johannes Welbl, Nelson F. Liu, and Matt Gardner. Crowdsourcing multiple choice science questions. In *Proceedings of the 3rd Workshop on Noisy User-generated Text*, 2017.

[23] Todor Mihaylov, Peter Clark, Tushar Khot, and Ashish Sabharwal. Can a suit of armor conduct electricity? a new dataset for open book question answering. In *Proceedings of the 2018 Conference on Empirical Methods in Natural Language Processing*, 2018.

[24] Denis Paperno, Germán Kruszewski, Angeliki Lazaridou, Quan Ngoc Pham, Raffaella Bernardi, Sandro Pezzelle, Marco Baroni, Gemma Boleda, and Raquel Fernández. The LAMBADA dataset: Word prediction requiring a broad discourse context. In *Proceedings of the 54th Annual Meeting of the Association for Computational Linguistics*, 2016.

[25] Pranav Rajpurkar, Jian Zhang, Konstantin Lopyrev, and Percy Liang. SQuAD: 100,000+ questions for machine comprehension of text. In *Proceedings of the 2016 Conference on Empirical Methods in Natural Language Processing*, 2016.